



Learning to Infer and Execute 3D Shape Programs

Yonglong Tian, Andrew Luo, Xingyuan Sun, Kevin Ellis, William T.
Freeman, Joshua B. Tenenbaum, Jiajun Wu

Motivation

- Unstructured “neural” representations



Under constrained

Not interpretable




What is the correct representation??

Idea of vision as “inverse graphics”


- Humans can easily understand visual input
 - Decomposition of a scene into objects
 - Decomposition of objects into components
 - Infer object relationships etc. etc.



More than just pixels in
2D image!



Taken from “Part-based modelling of compound scenes from images”, van den Hengel, Anton, et al.



OR

```
Draw("Top", "Circle", pos0, shape_param0)
Draw("Leg", "Cub", pos1, shape_param1)
Draw("Leg", "Cub", pos2, shape_param2)
Draw("Leg", "Cub", pos3, shape_param3)
Draw("Leg", "Cub", pos4, shape_param4)
Draw("Layer", "Rec", pos5, shape_param5)
Draw("Layer", "Rec", pos6, shape_param6)
```

OR



```
Draw("Top", "Circle", position, geometry)
```

```
for(i < 2, "translation", a)
```

```
for(j < 2, "translation", b)
```

```
Draw("Leg", "Cub", position + i*a + j*b, geometry)
```

```
for(i < 2, "translation", c)
```

```
Draw("Layer", "Rec", position + i*c, geometry)
```



Our model

```
Draw("Top", "Circle", position, geometry)
```

```
for(i < 2, "translation", a)
```

```
  for(j < 2, "translation", b)
```

```
    Draw("Leg", "Cub", position + i*a + j*b, geometry)
```

```
for(i < 2, "translation", c)
```

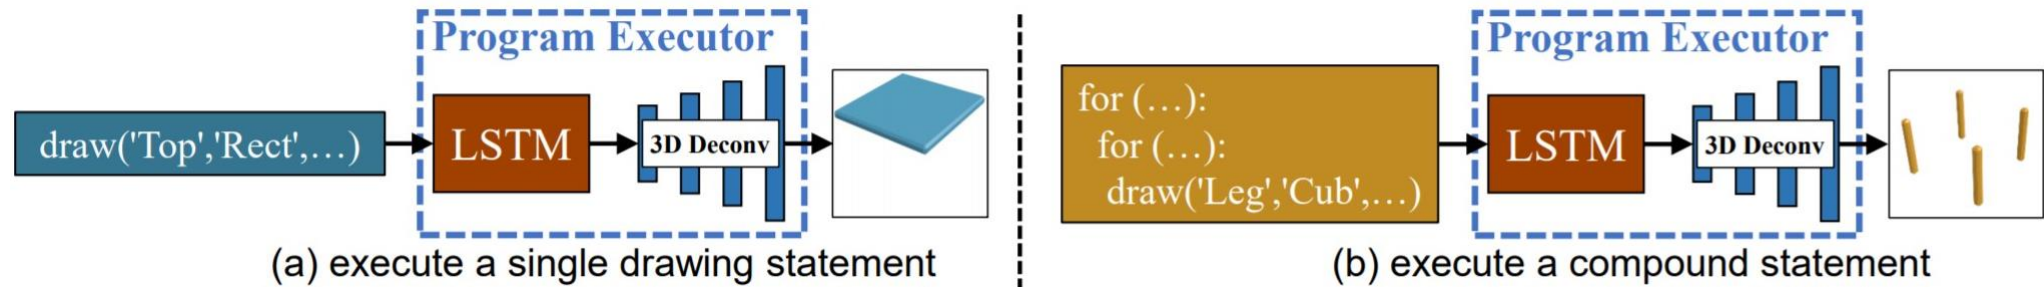
```
  Draw("Layer", "Rec", position + i*c, geometry)
```

Challenges

- No dataset of sufficiently high quality exists for part based decomposition of 3D objects
 - Roll your own “synthetic” dataset
- We have templates for chairs, tables
 - Can be fuzzed for diversity
 - Hard codes loops for translation & rotation (think office chair base)
 - Possibly disjoint from “real” data

Model Components:

1. Program rendering module (program executor)
Takes in a block of code, uses RNN + 3d deconv to output voxels



Trained purely on synthetic data

Model Components:

2. Program output LSTM

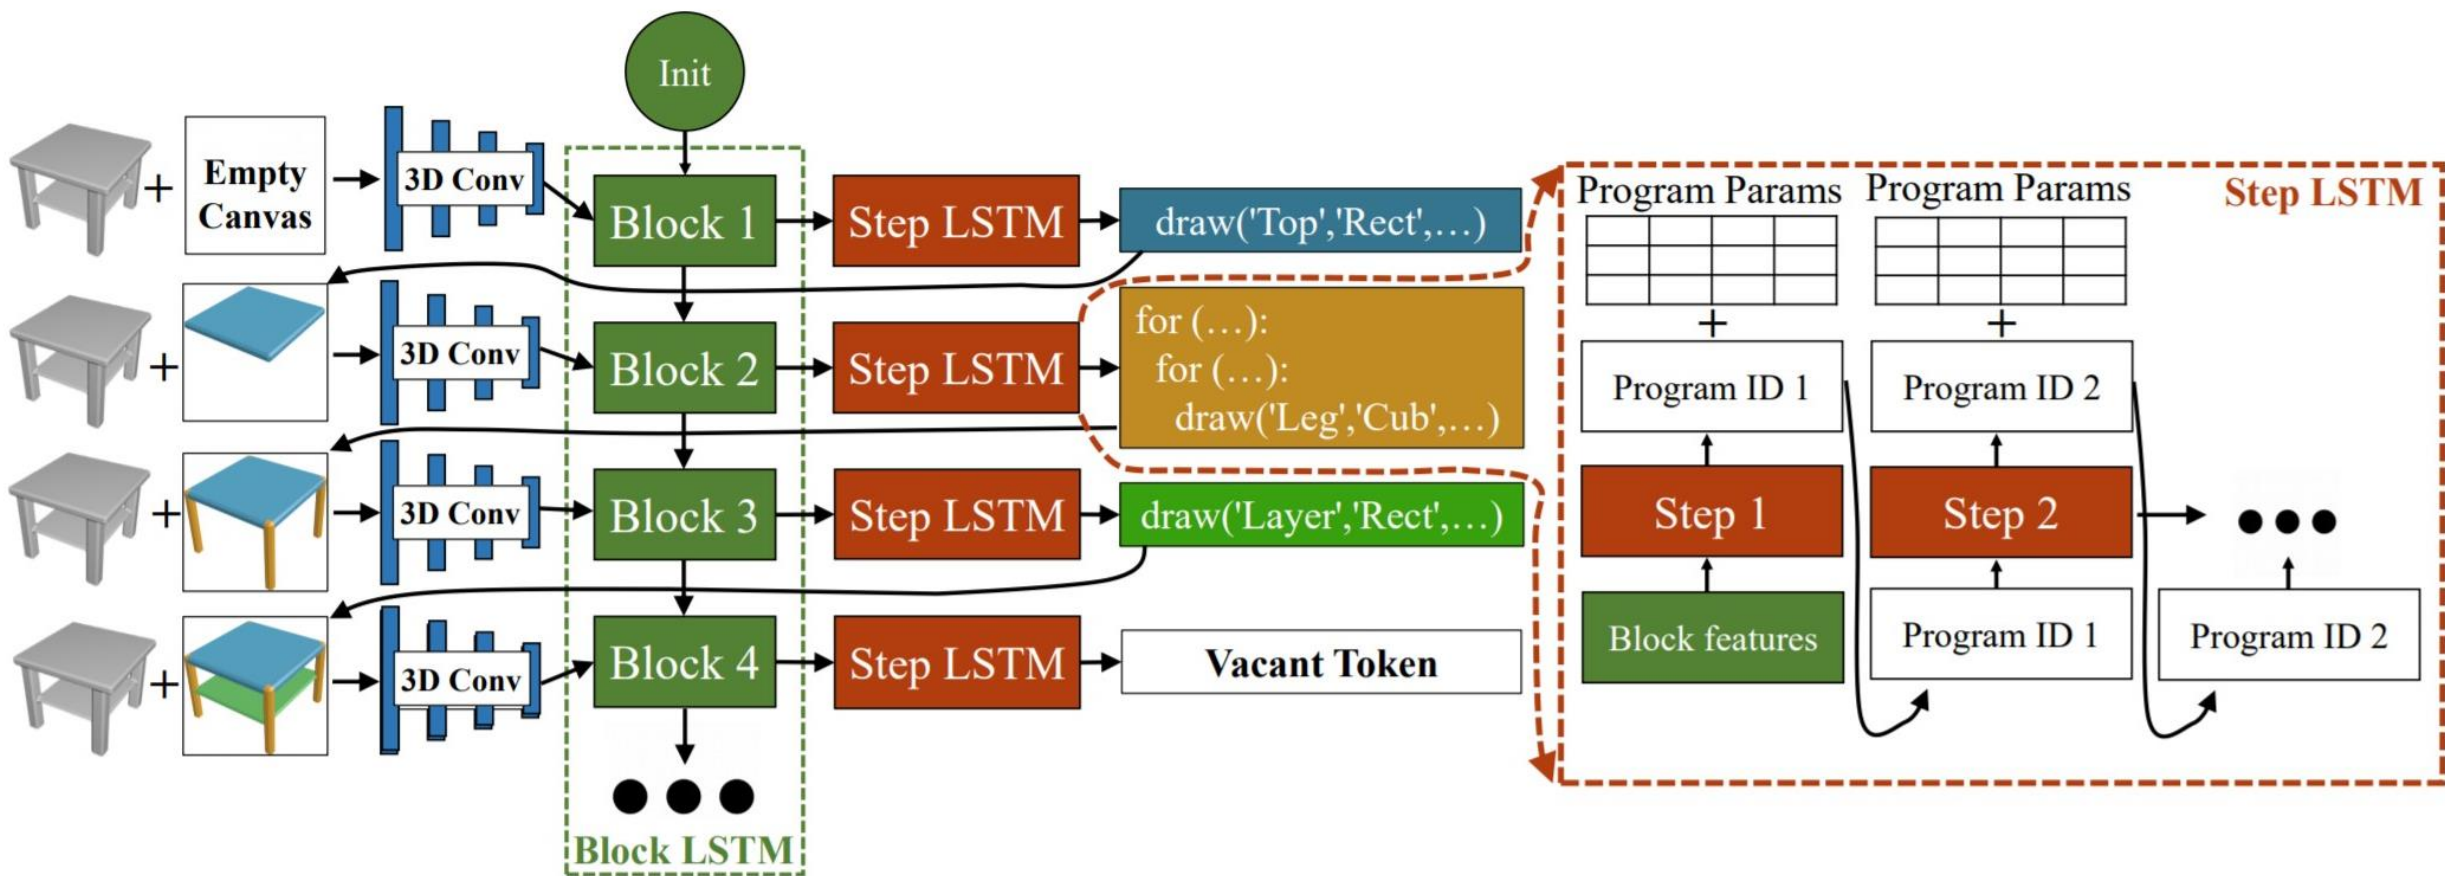
Takes in:

- input voxels
- canvas (to keep track of already rendered voxels)

Actually two LSTMs, one of them inside the other one

First LSTM identifies blocks of voxels that belong together

Second LSTM takes the features produced, and outputs programs



Why do we even need the executor

- Program type can be trained via cross-entropy
- Program parameters can be trained via L2
- To visualize we can just run the programs right??

Fine tuning (Guided adaptation)

- Synthetic data might be different from real data
- Remember, program executor is a neural network
- Neural network = differentiable

For a given ground truth object:

1. Generate programs
2. Execute, and backprop loss to program generator to tune
3. Repeat to get better at a given ground truth object

Results:

Models	IoU \uparrow		CD \downarrow		EMD \downarrow	
	table	chair	table	chair	table	chair
CSGNet-original	0.111	0.154	0.216	0.175	0.205	0.177
Tulsiani et al. (2017)	0.357	0.406	0.083	0.079	0.073	0.072
CSGNet-augmented	0.406	0.365	0.072	0.077	0.069	0.076
Nearest Neighbour	0.445	0.389	0.083	0.084	0.084	0.084
Shape Programs w/o GA	0.487	0.422	0.067	0.072	0.063	0.072
Shape Programs	0.591	0.516	0.058	0.063	0.056	0.060

Models	Stable (%)		Conn. (%)		Stable & Conn. (%)	
	table	chair	table	chair	table	chair
Tulsiani et al. (2017)	36.7	31.3	37.1	68.9	15.4	19.6
Shape Programs w/o GA	94.7	95.1	76.6	54.2	73.7	51.6
Shape Programs	97.0	96.5	78.4	68.5	77.0	66.0
Ground Truth	98.9	97.6	98.8	97.8	97.7	95.5

Before and after fine tuning



Reconstruction before adaption

```
draw('Top', 'Cir', P=(4,0,0), G=(1,7))
draw('Support', 'Cyl', P=(-9,0,0), G=(15,3))
for(i<4, Rot( $\theta_{rot}=90$ , ax=(-9,1,0)))
  draw('Base', 'Line', P=(-9,1,0),
      G=(-9,-6,-5),  $\theta_{rot} \times i$ , ax)

draw('Layer', 'Rect', P=(-3,0,0), G=(2,4,6))
```



Input

Reconstruction after adaption



```
draw('Top', 'Cir', (P=(0,0,0), G=(2,6)))
draw('Support', 'Cyl', P=(-11,0,0), G=(13,1))

for(i<5, 'Rot',  $\theta_{rot}=72$ , ax=(-10,0,0))
  draw('Base', 'Line', P=(-10,0,0),
      G=(-11,-6,-3),  $\theta_{rot} \times i$ , ax)

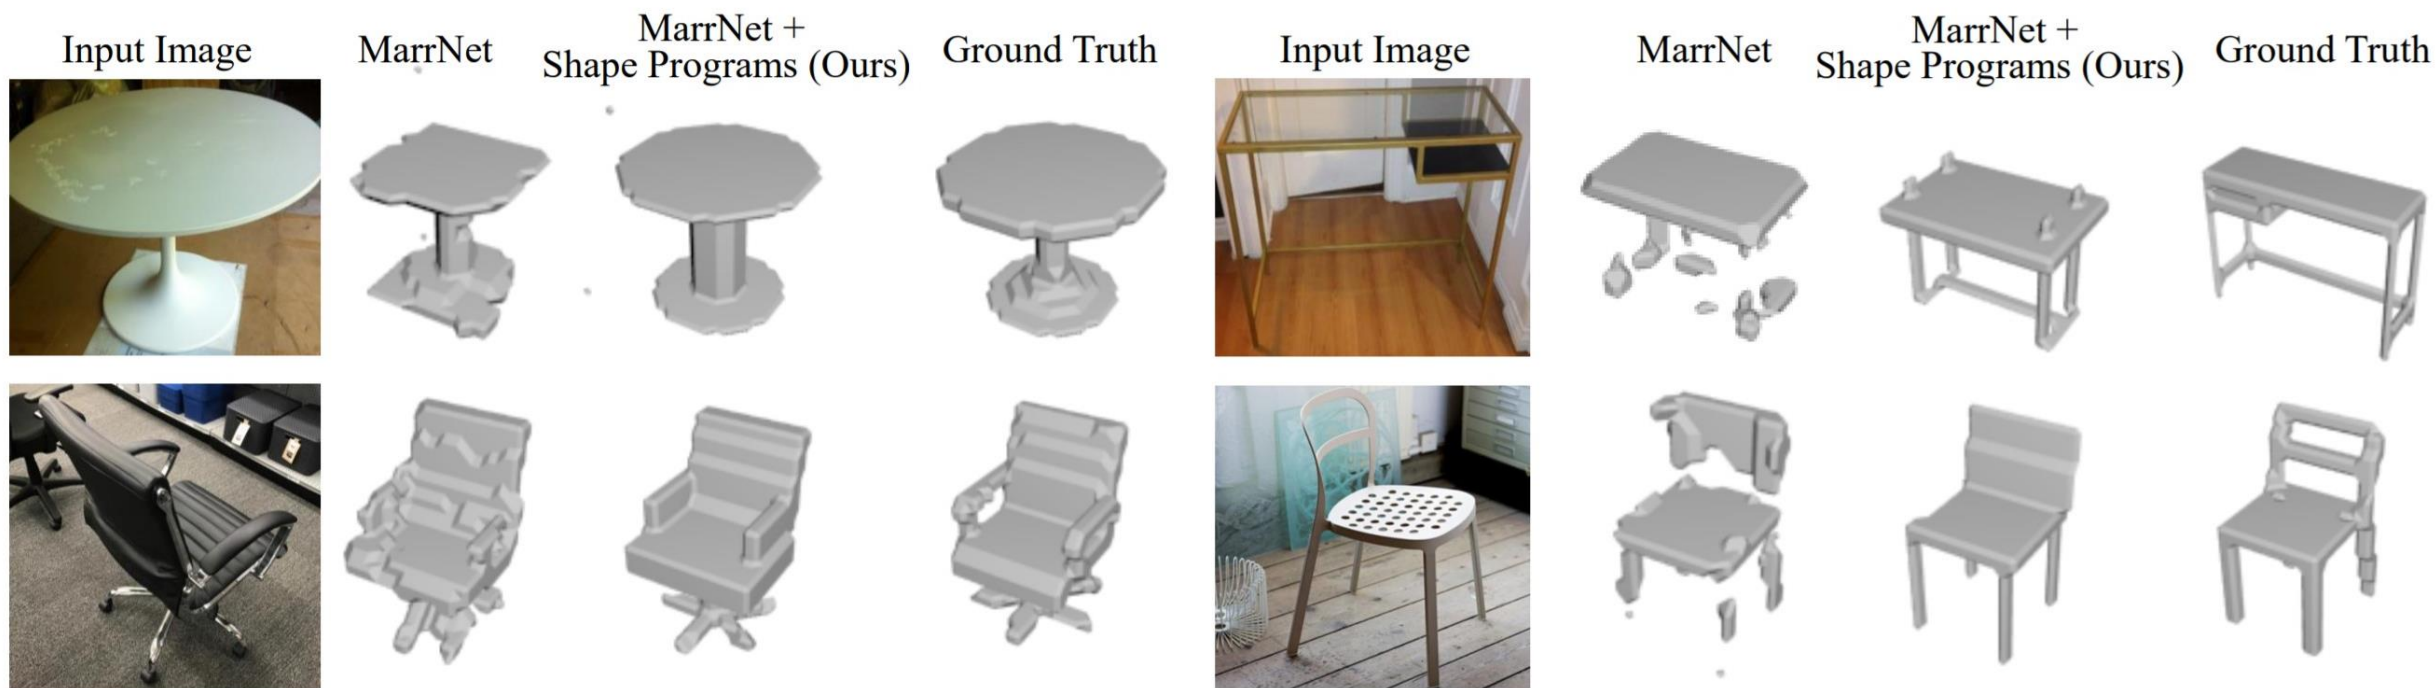
draw('TiltBack', 'Cub', P=(3,2,-5), G=(8,2,9,7))

for(i<2, 'Trans', u1=(0,0,11))
  for(j<2, 'Trans', u2=(0,4,0))
    draw('ChairBeam', 'Cub', P=(2,-4,-6)
        + (ju2) + (iu1), G=(3,1,2))

for(i<2, 'Trans', u=(0,0,10))
  draw('HoriBar', 'Cub', P=(4,-4,-6)
      + (iu), G=(1,5,2))
```


(a)

Constrained representation on reconstruction



Project 2


Object relationships can be represented with graph




Taken from “Image Generation from Scene Graphs” by Johnson et al.



Previous approaches

- Image generation from scene graph (Johnson) – CVPR 2018





A person riding a wave and a board by the water with sky above.




(1)




A boy standing on grass looking at a kite and the sky with the field under a mountain



(2)



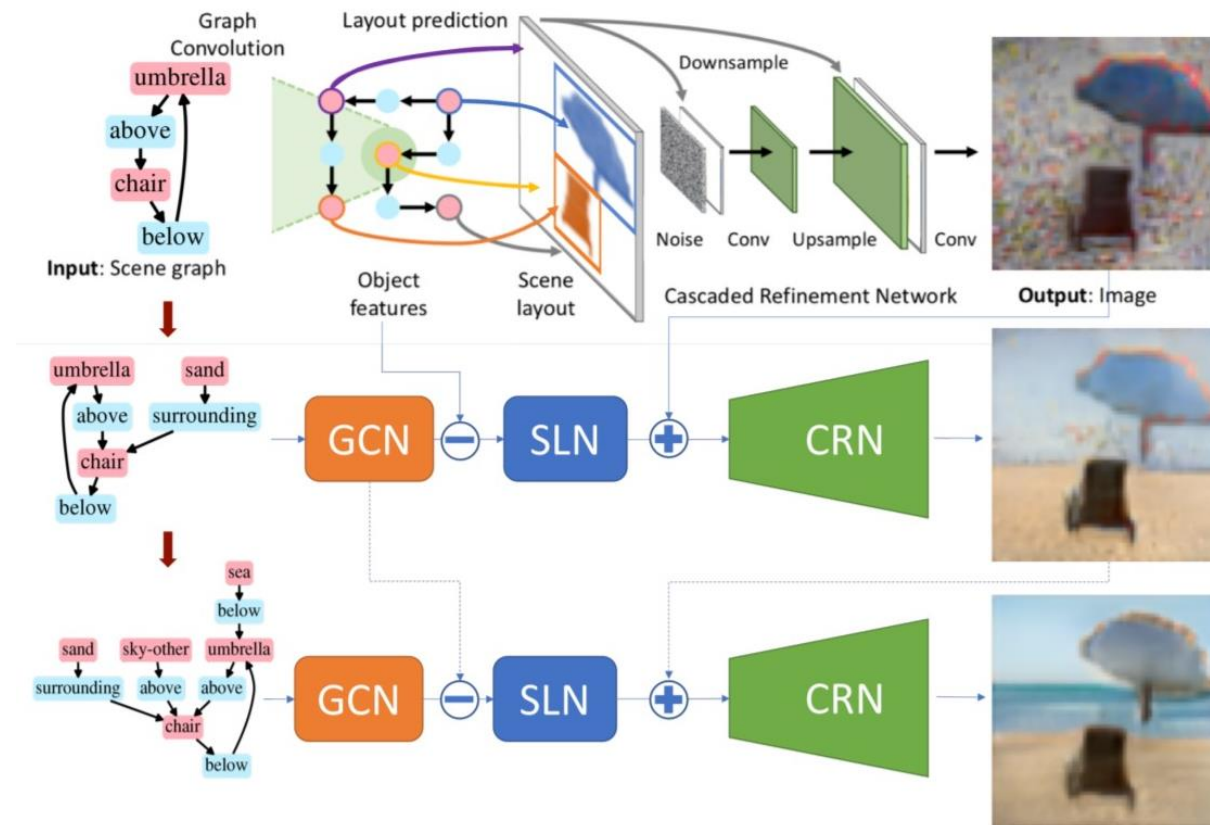
Two busses, one behind the other and a tree behind the second; both busses have windshields.



(3)

Previous approaches

- Interactive Image Generation Using Scene Graphs – ICLR 2019 workshop



Bonus Paper:

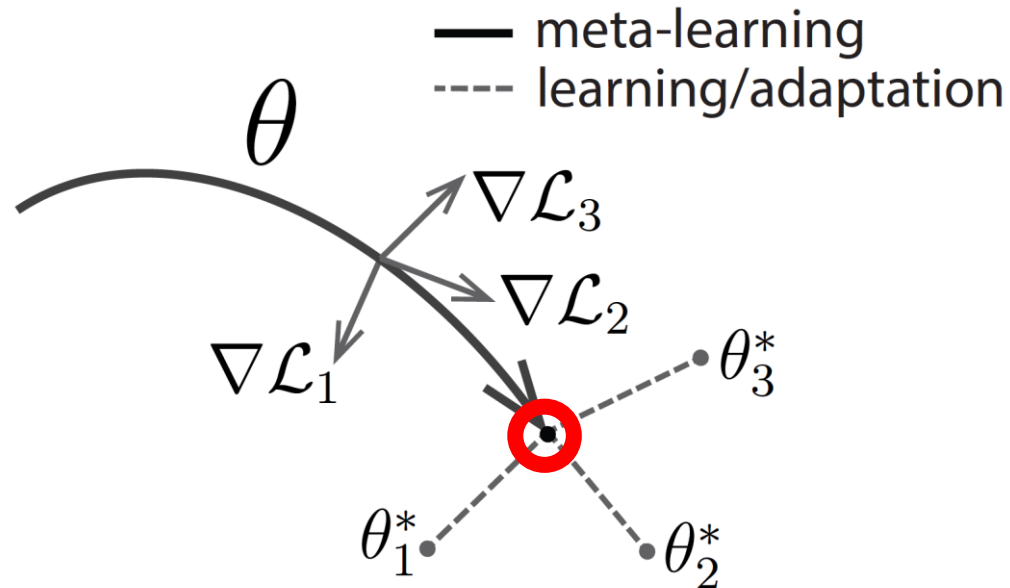
MAML: Model-Agnostic Meta-learning

Meta Learning

- Learning to learn
 - Optimize model for fast adaptation
- Closely related to few shot learning
- Should be able to adapt to new tasks easily
- Reuse learned concepts
- See “Human-level concept learning through probabilistic program induction” by Tenenbaum et al.

MAML

- Optimization weights for better generalization
- Applies to models trained via gradient descent
- Requires higher order gradients support
 - Pytorch (broke on multi-GPU since 1.0, fixed in 1.2), Tensorflow
 - Mxnet work in progress, probably within the next 2 versions (just a guess)



Approach

- Explicitly optimize θ such that the following is minimized:

$$\sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})})$$

- Note: $p(\mathcal{T})$: distribution over tasks

α : Task learning rate (think of this as the fine tuning amount)

Approach

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

1: randomly initialize θ

2: **while** not done **do**

3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$

4: **for all** \mathcal{T}_i **do**


5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples

6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$

7: **end for**

8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$

9: **end while**



Bonus-bonus paper

- Reptile: A Scalable Meta- Learning Algorithm
 - Faster
 - First order approximation